



Introduction to

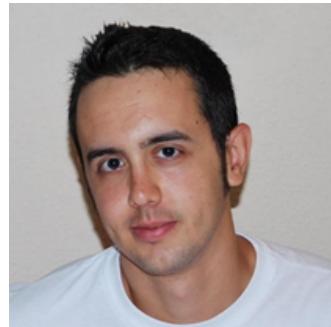


Javier Santos

May, 11 - 2015



About me



Javier Santos *@jpaniego*

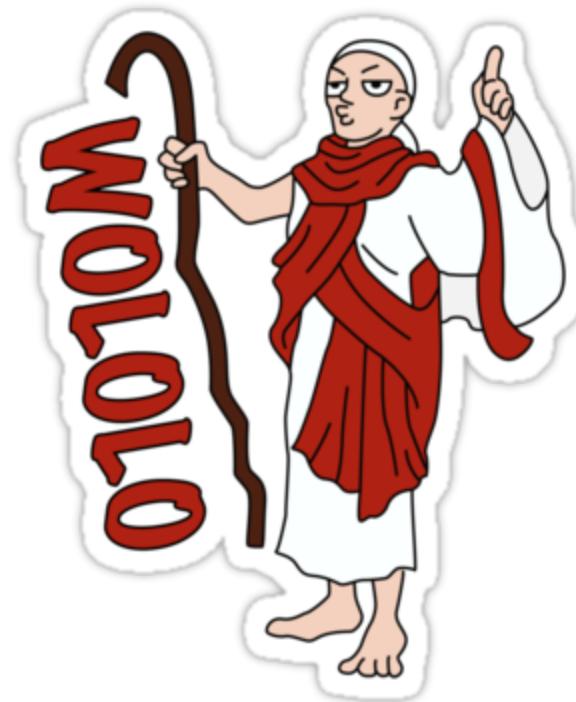
«Hay dos formas de programar sin errores; solo la tercera funciona»

Alan J Perlis



Disclaimer I

- I'm no Spark evangelizer





Disclaimer II

- I'm no expert





Disclaimer III

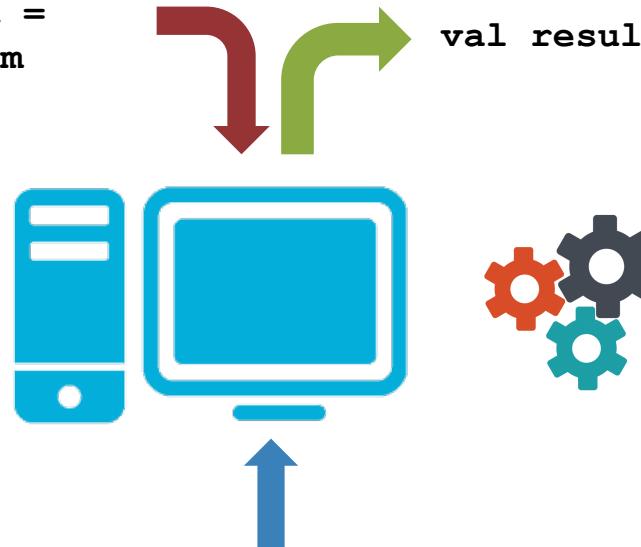
- I'm not paying beers





What's Spark?

```
val action: Seq[Int] => Int =  
  (seq: Seq[Int]) => seq.sum
```

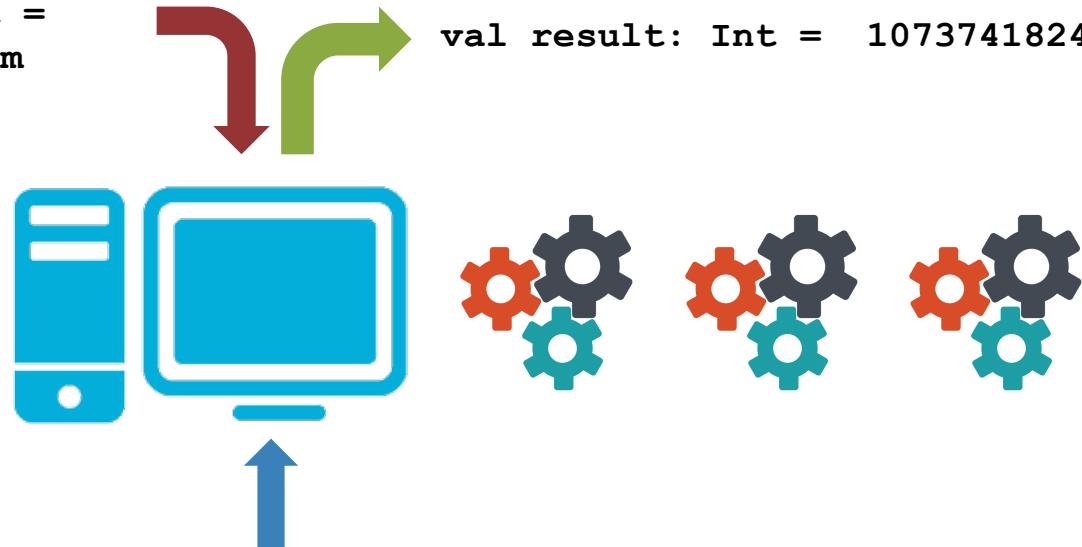


```
val mySeq: Seq[Int] = 1 to 5
```



What's Spark?

```
val action: Seq[Int] => Int =  
  (seq: Seq[Int]) => seq.sum
```



```
val mySeq: Seq[Int] = 1 to Int.MaxValue
```



What's Spark?

Computing problems:

- Time: Waiting for action results on tons of records
- Space: How do we allocate TB of data?

1/20/04 at 4:12pm **THREAD STARTER**

post #1 of 26

SpcMs ▾

Joined: Jun 2003

Posts: 379

 offline

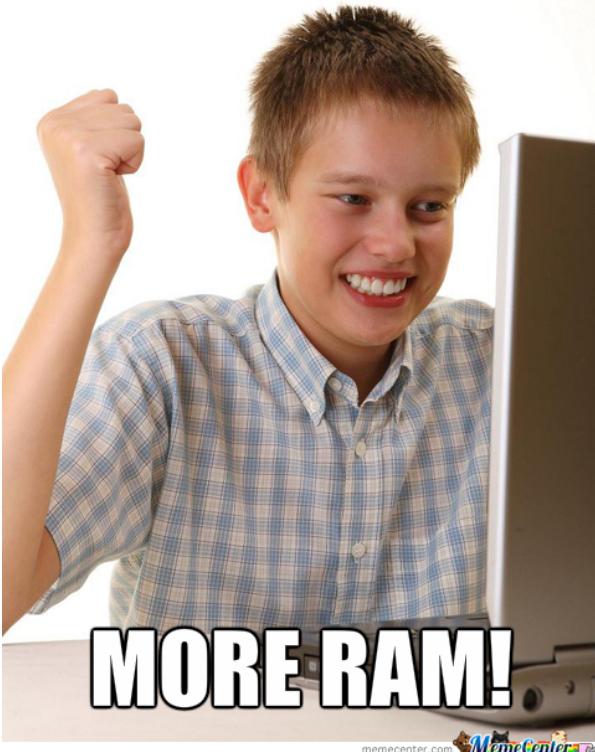
Any good sites that I can download some RAM from? My computer tells me i'm low on memory. Thanks in advance

It's Better To Be Hated For What You Are Than To Be Loved For What You Are Not

downloadmoreram.com

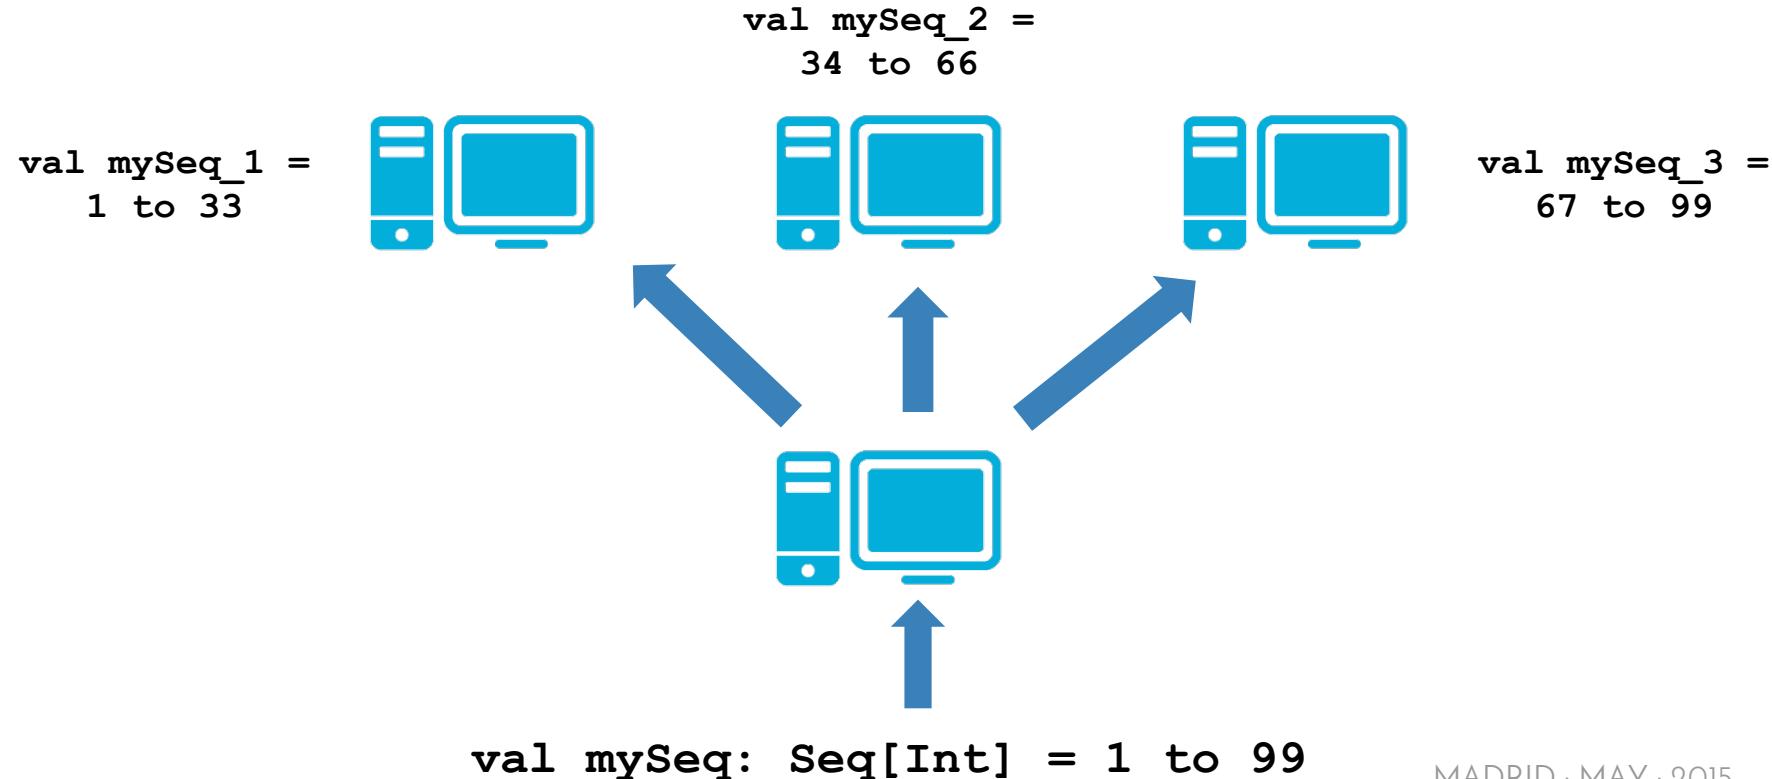


FINALLY I DOWNLOADED



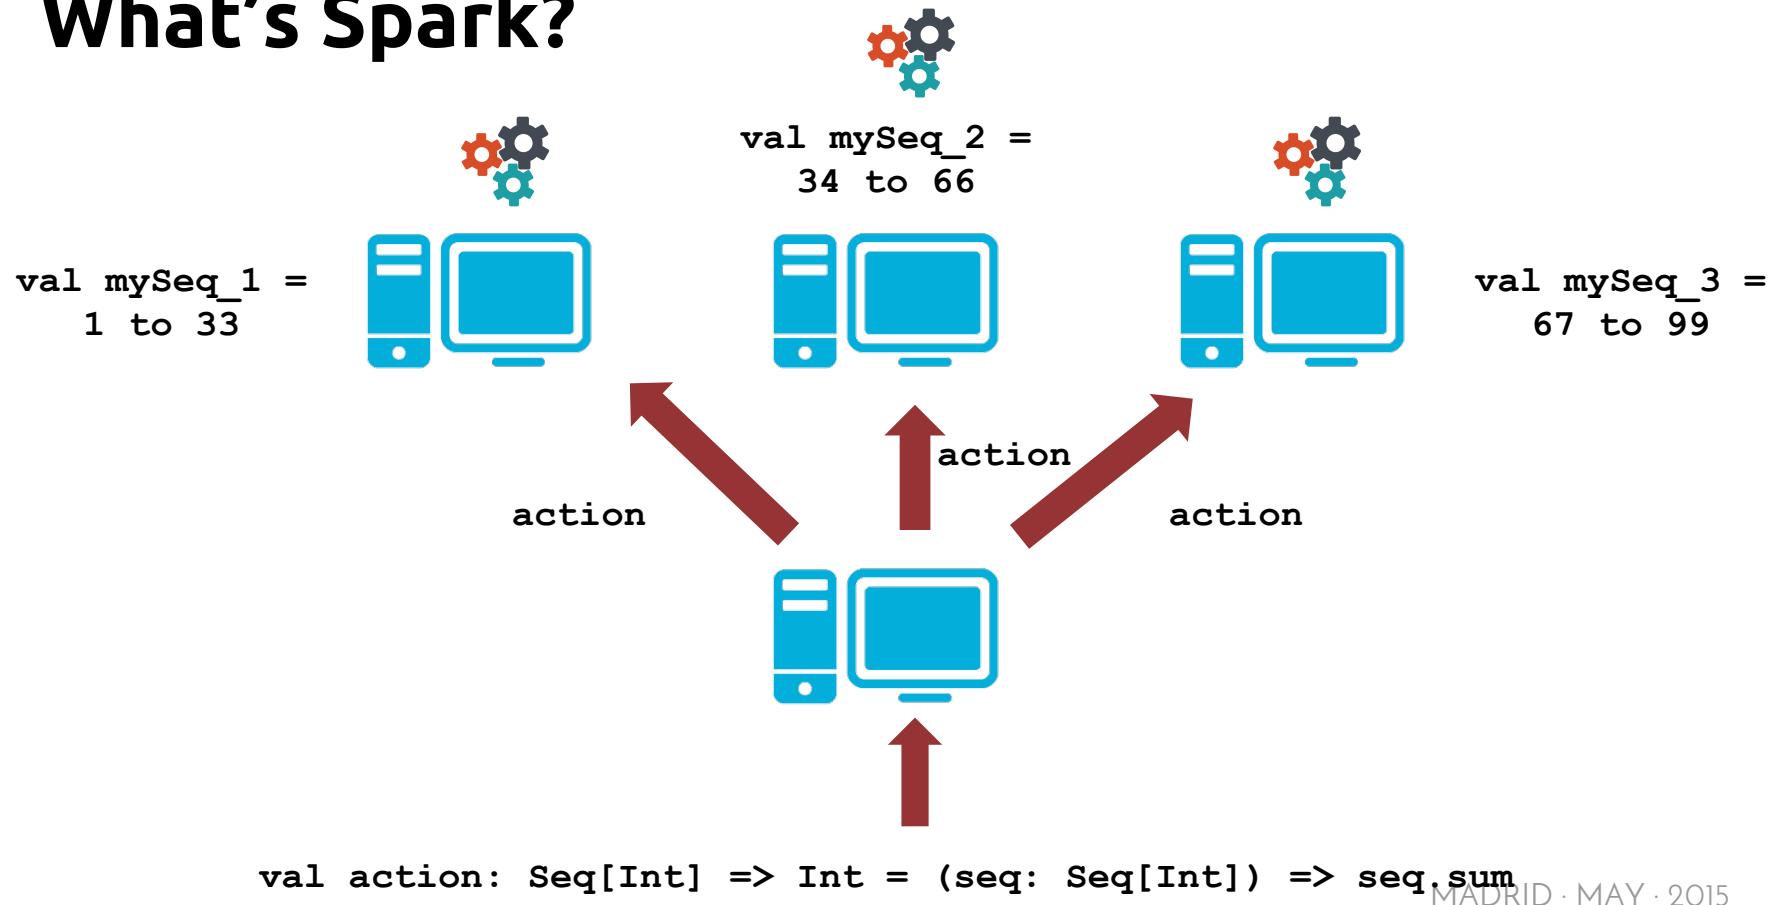


What's Spark?



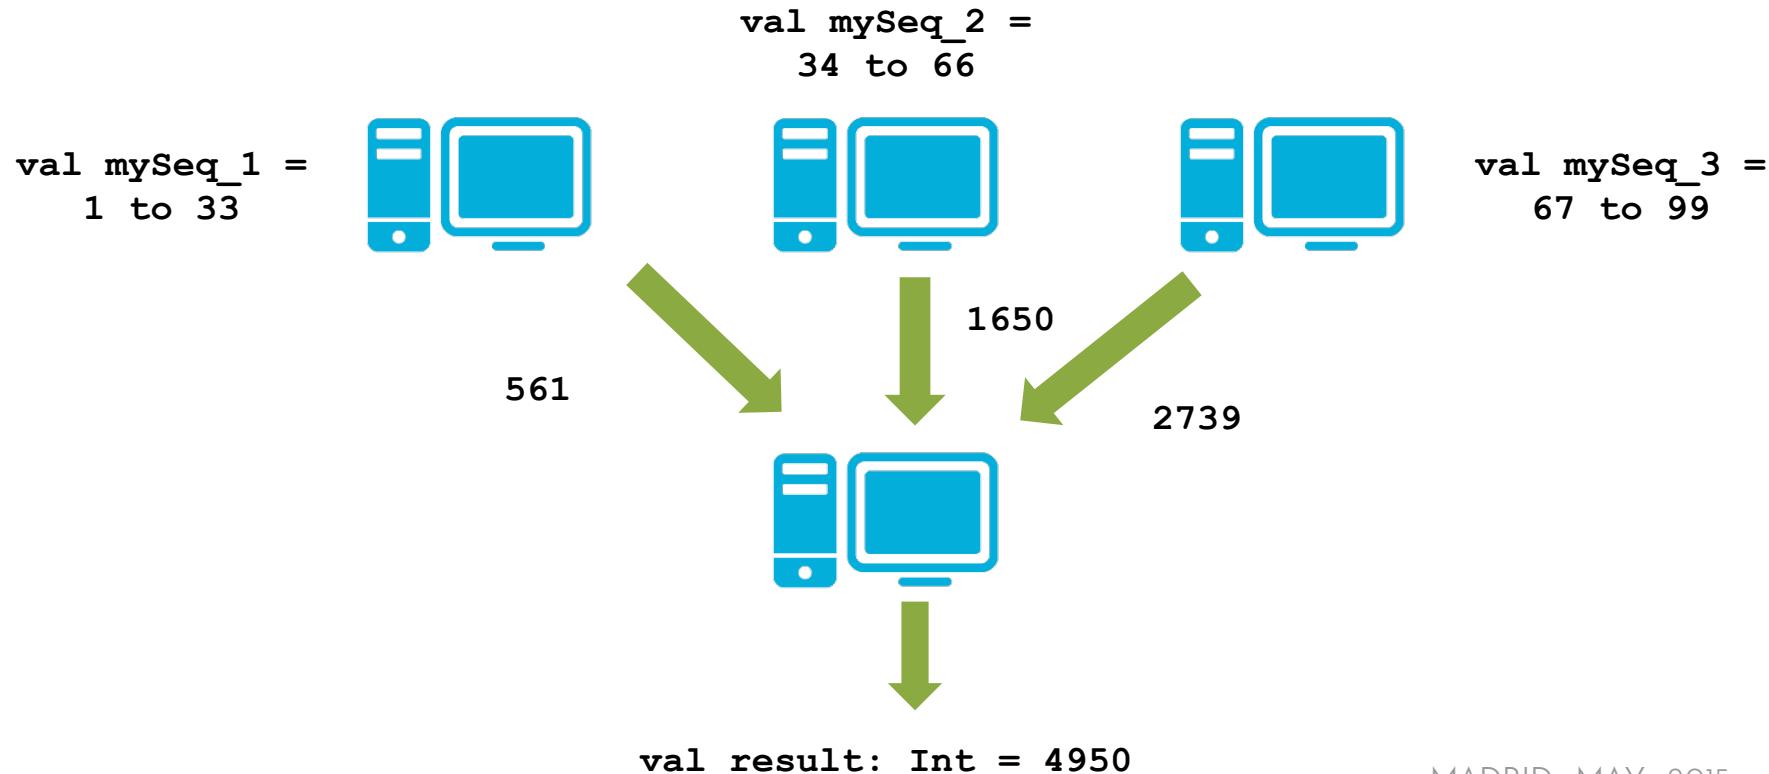


What's Spark?





What's Spark?





What's Spark?

- General engine for large-scale data processing.
- “In-memory” data allocation.
- Target: data and computing parallelism.
- Scala, Python and Java APIs.
- First developed by [AMPLab](#) in UC Berkeley
- Support by [Databricks](#)
- Latest stable version 1.3.1



Modules

Spark
SQL

Spark
Streaming

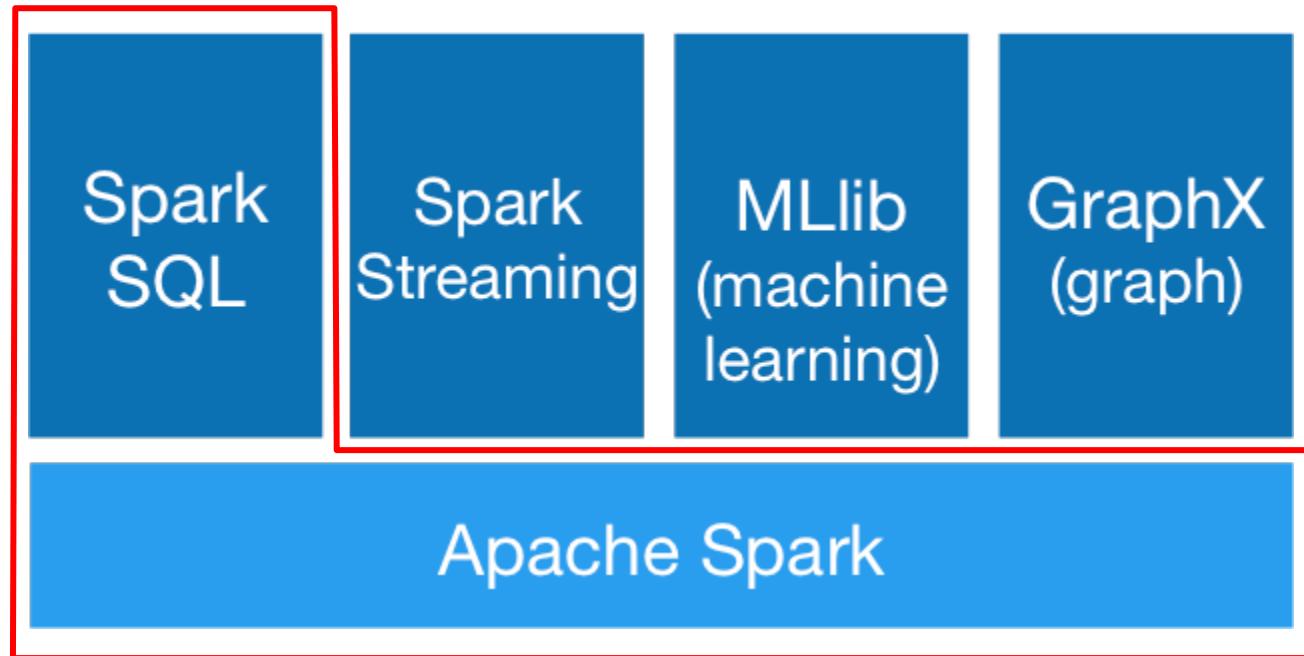
MLlib
(machine
learning)

GraphX
(graph)

Apache Spark



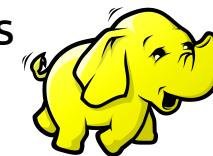
Modules





MapReduce model

- Linear calculus method used for supporting parallel-computing and big data sets distribution among computing groups
- 80s. Google. PageRank. Big arrays computation.
- Based on two main functional programming methods: Map and reduce.
- Hadoop
 - One of the first open source implementations
 - Most of contributions made by Yahoo





MapReduce model

- Phases

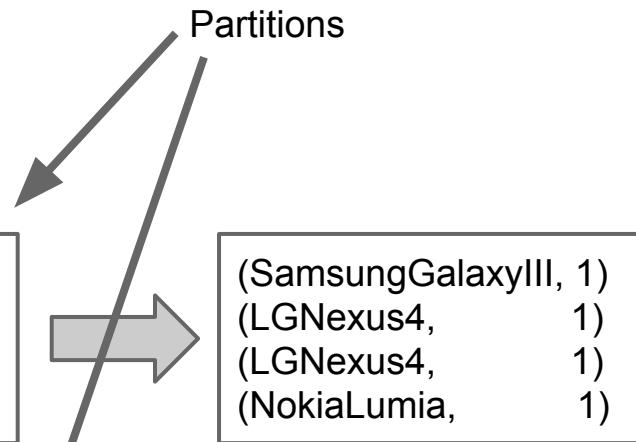




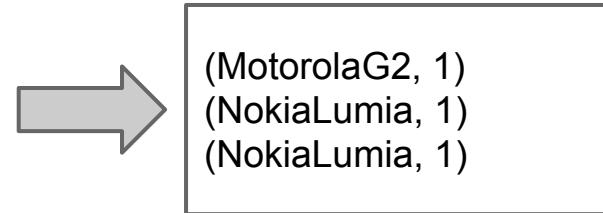
MapReduce model

- Example (Map)

```
"123;SamsungGalaxyIII;14:05;300;34612345678"  
"124;LGNexus4;16:05;121;+34613455678"  
"126;LGNexus4;12:05;23;+3463624678"  
"131;NokiaLumia;14:05;300;+34613246778"
```



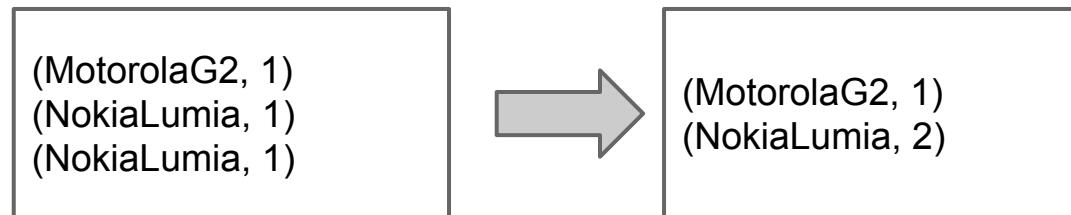
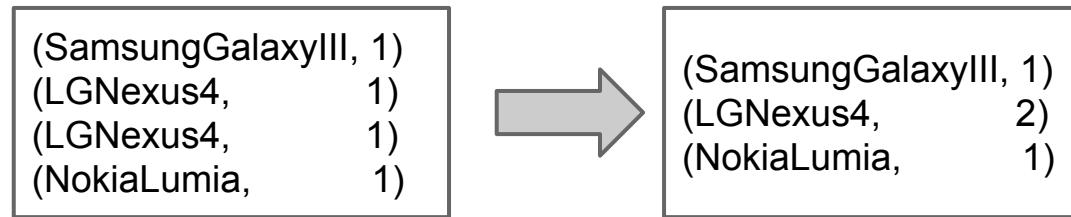
```
"125;MotorolaG2;14:05;300;+34612345678"  
"127;NokiaLumia;14:05;300;+34612345678"  
"130;NokiaLumia;14:05;300;+34612345678"
```





MapReduce model

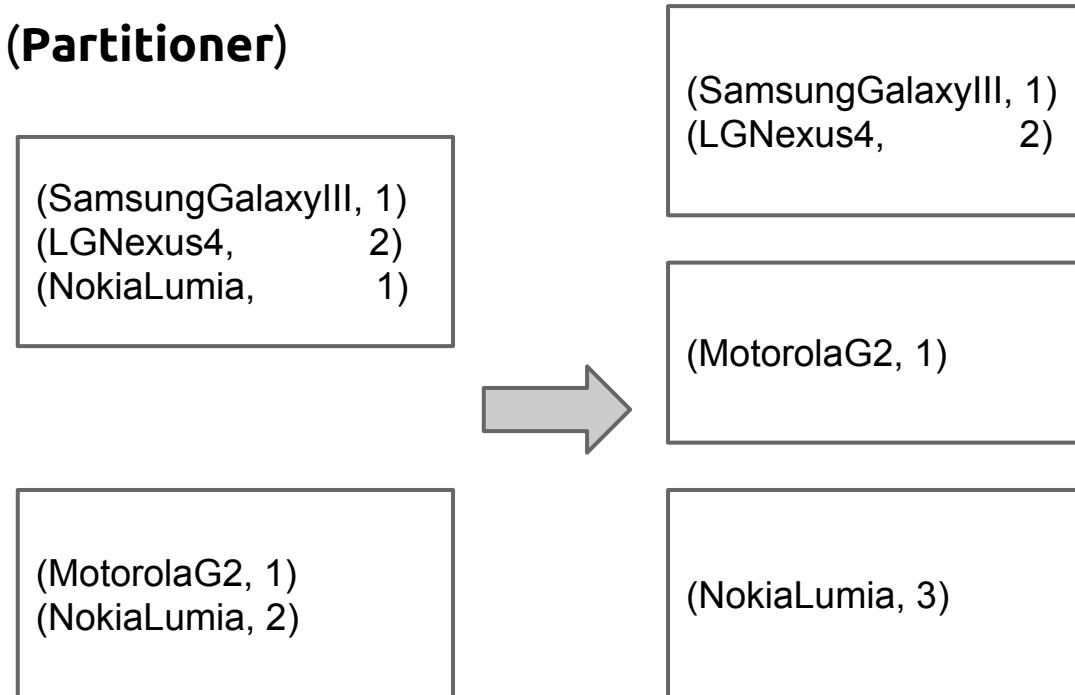
- Example (**Combiner - local aggregator**)





MapReduce model

- Example (**Partitioner**)





MapReduce model

- Example (**Reduce**)

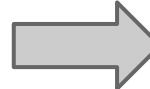
(SamsungGalaxyIII, 1)
(LGNexus4, 2)

(MotorolaG2, 1)

(NokiaLumia, 3)

ReduceFunction

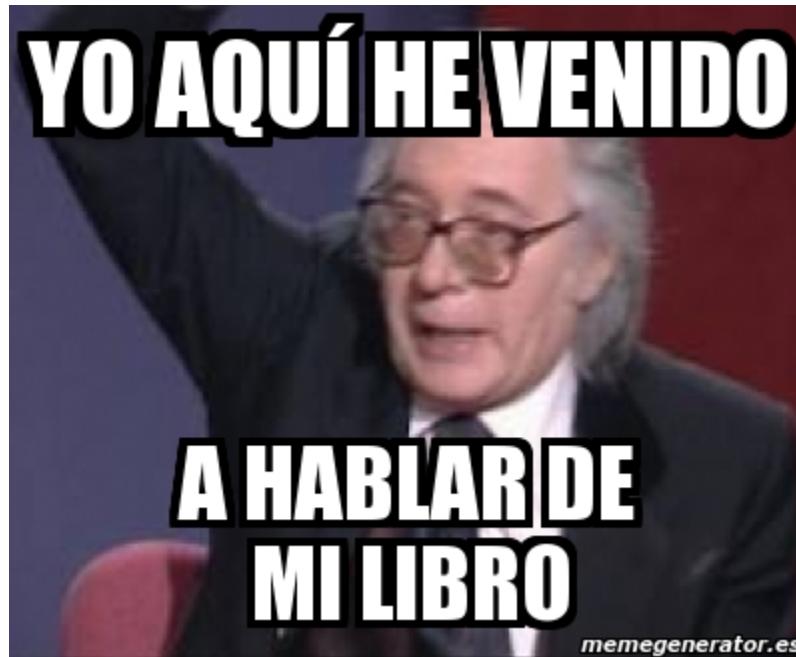
```
f(t1: (String, Int), t2: (String, Int)) =  
  if (t1._2 > t2._2) t1 else t2
```



(NokiaLumia, 3)

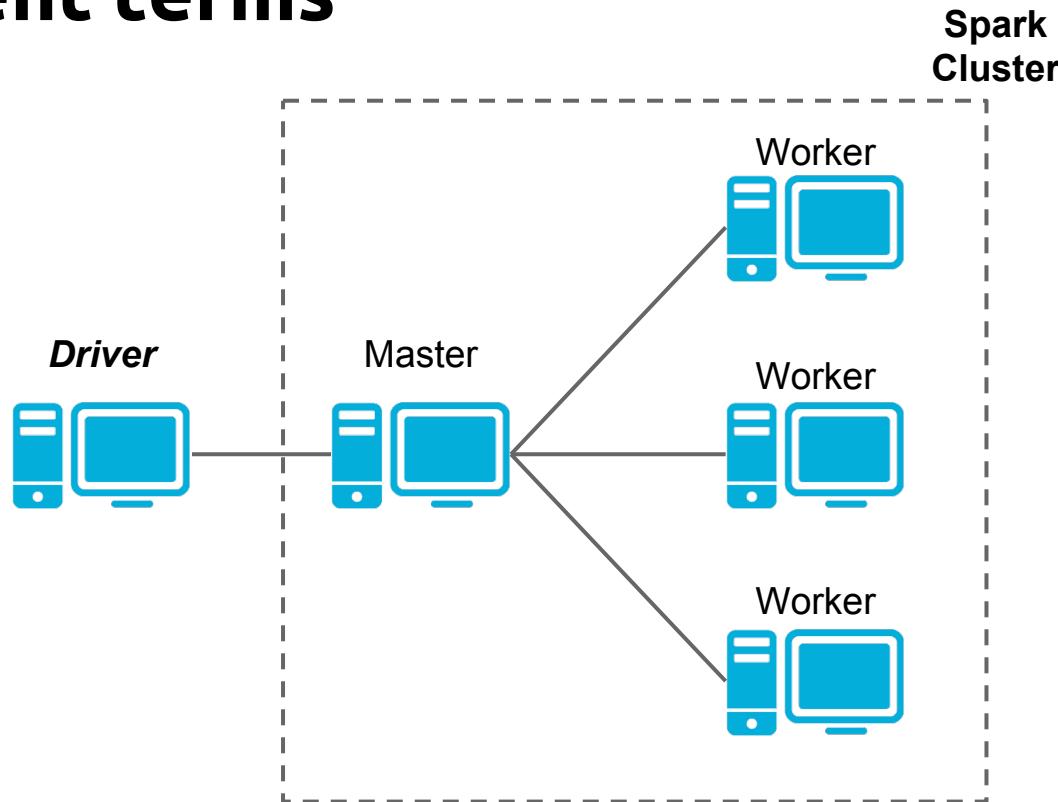


So...what about Spark?





Deployment terms



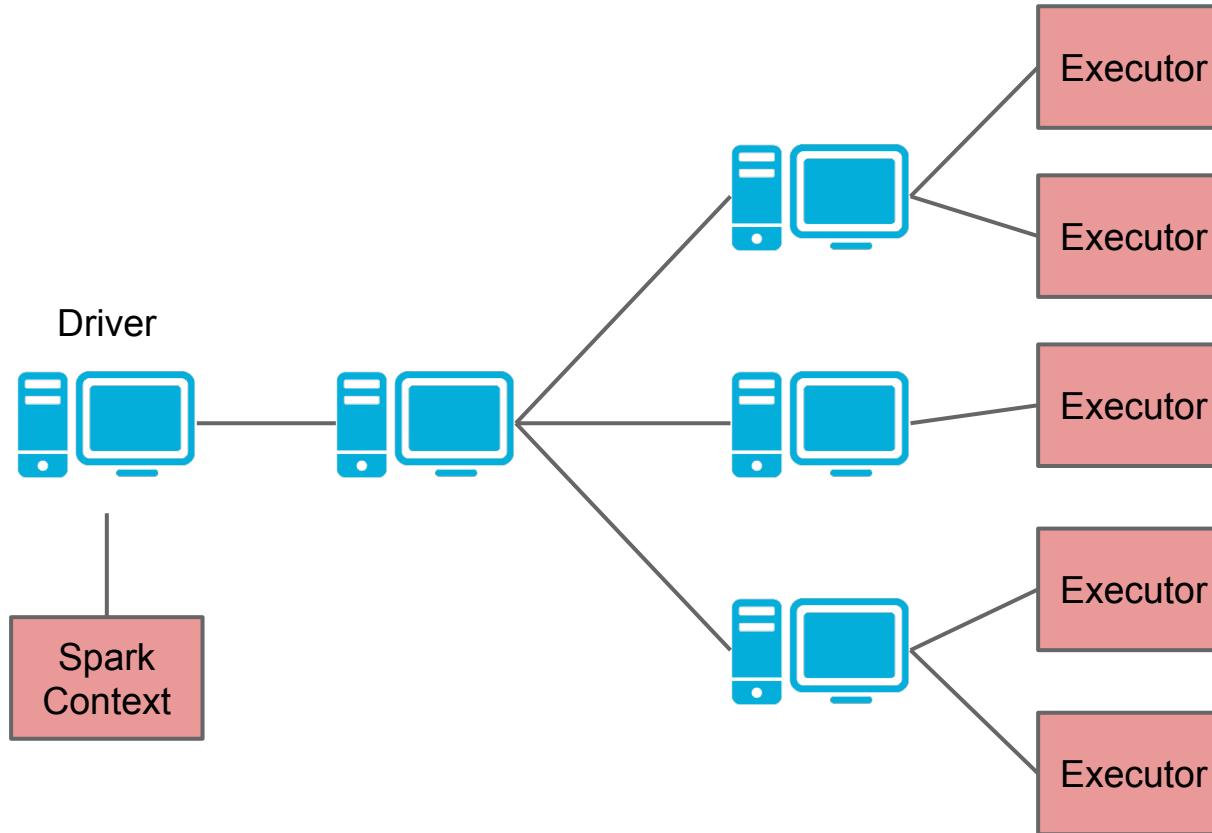


Deployment types

- Local
 - master="local [N]" (N=Amount of cores)
 - Spark master is launched in the same process.
 - Workers are launched in the same process.
- Standalone
 - master="spark://master-url"
 - Spark master is launched in a cluster machine.
 - Submit JAR to worker nodes.
 - Resource managers: YARN, Mesos



Execution terms



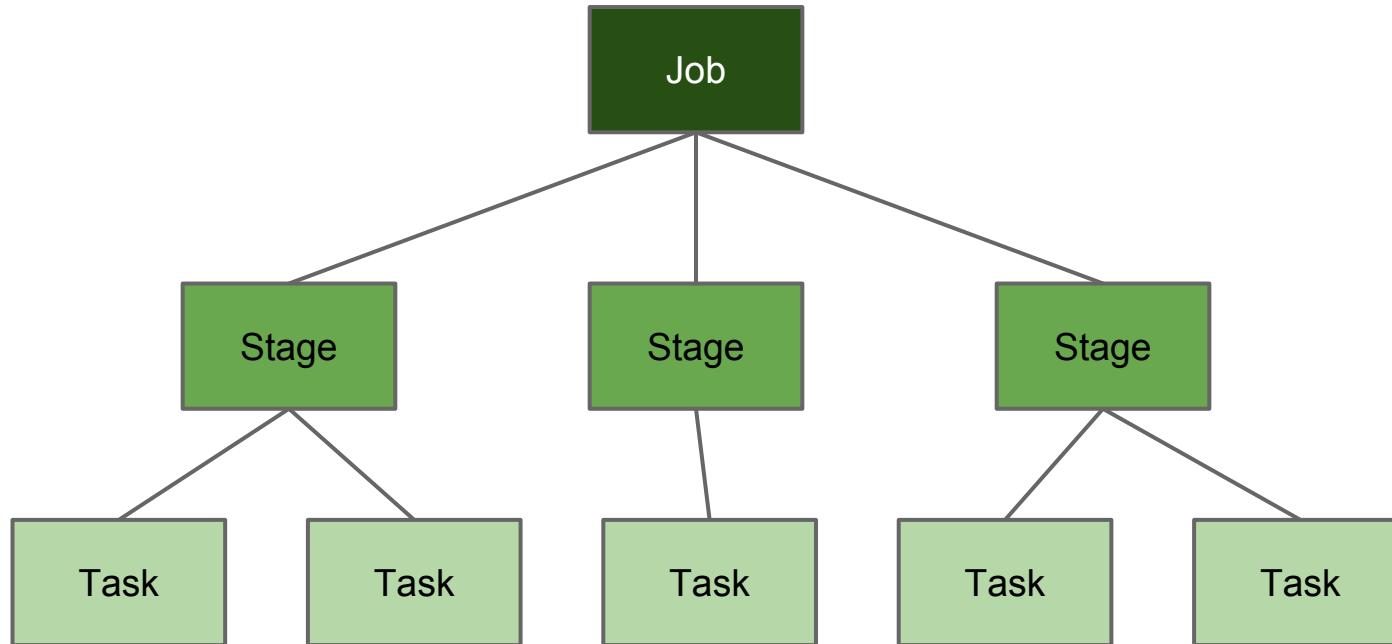


Execution terms : SparkContext & Executor

- **SparkContext**
 - Spark cluster connection
 - Necessary for SQLContext, StreamingContext, ...
 - Isolated: Two Spark Contexts cannot exchange data (without external help).
- **Executor**
 - Individual execution unit
 - 1 core ~ 2 executor
 - Each application has its own.



Job terms





Modules

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Apache Spark



RDD

- **Resilient Distributed Dataset.** Basic abstraction in Spark
- Think of it like a huge collection partitioned and distributed in different machines.
- Lazy evaluated
- Basically composed of
 - A list of partitions
 - A function for computing each split
 - A list of dependencies on other RDDs



RDD

- Basic example:

```
val myRdd: RDD[String] = sc.parallelize(  
    "this is a sample string".split(" ").toList)
```

```
val anotherRdd: RDD[(Int, String)] = sc.parallelize(  
    List(  
        0 -> "A value for key '0' ",  
        1 -> "Another value"))
```



RDD

- An RDD may be created from

- A file / set of files:

```
sc.textFile("myFile")
```

```
sc.textFile("file1,file2")
```

- A bunch of memory-storaged data:

```
sc.parallelize(List(1,2,3))
```

Another RDD:

```
myRdd.map(_.toString)
```

- Another RDD



RDD - Partitions

- Partition : RDD chunk
- A worker node may contain 1 or more partitions of some RDD.
- It's important to choose a best-performance partitioning.
- Repartitioning
 - It implies shuffling all RDD data among worker nodes
 - There will surely be tons of network traffic among all worker nodes!



RDD - Lineage

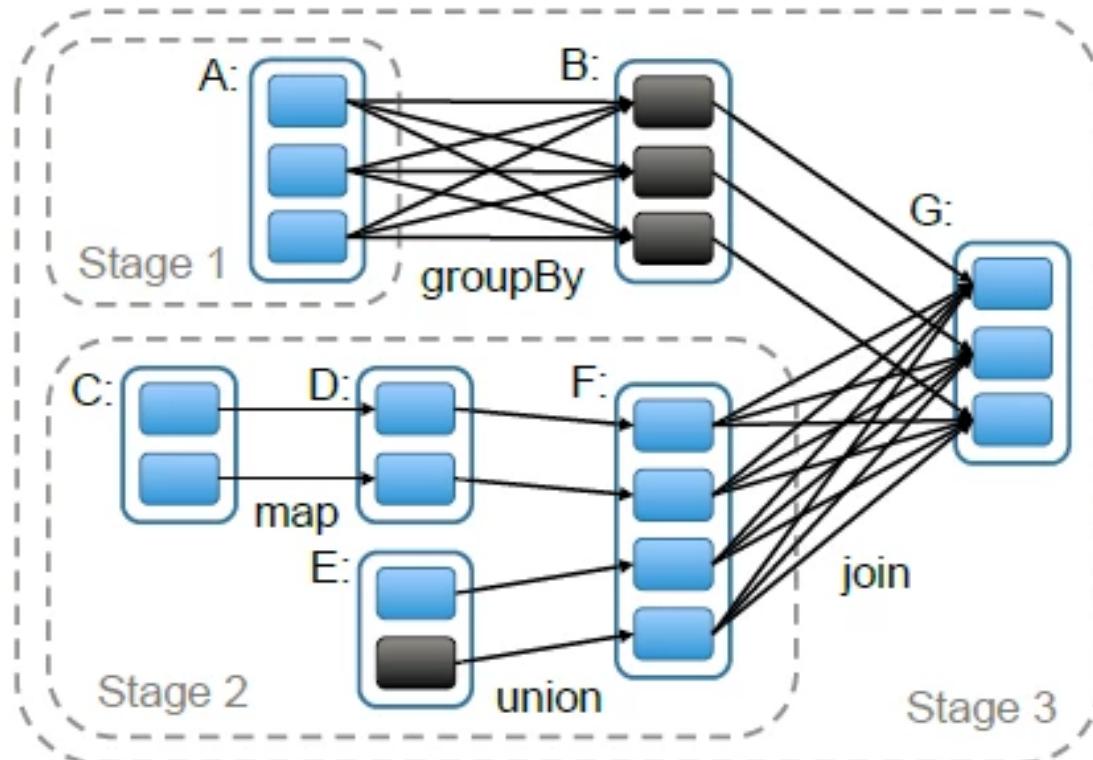
- An RDD may be built from another one.
- Base RDD : Has no parent
- Ex:

```
val base: RDD[Int] = sc.parallelize(List(1,2,3,4))  
val even: RDD[Int] = base.filter(_%2==0)
```

- DAG (**D**irected **A**cyclic **G**raph) : Generalization of MapReduce model.
Manages RDD lineages (source, mutations, ...).



RDD - Lineage





RDD - Transformations

- It applies a change to some RDD, returning a new one.
- It's not immediately evaluated.
- Think of it like a “Call-by-name” function. λ
- Adds a new node in Lineage DAG.



RDD - Transformations

- Most frequently used:

- **map**

```
val myRDD: RDD[Int] = sc.parallelize(List(1,2,3,4))
```

```
val myStringRDD: RDD[String] = myRDD.map(_.toString)
```

- **flatMap**

```
val myRDD: RDD[Int] =
```

```
sc.parallelize(List(1,null,3,null))
```

```
val myNotNullInts: RDD[Int] =
```

```
myRdd.flatMap(n => Option(n))
```



RDD - Transformations

- Most frequently used:

- **filter**

```
val myRDD: RDD[Int] = sc.parallelize(List(1,2,3,4))
```

```
val myStrinRDD: RDD[Int] = myRDD.filter(_%2==0)
```

- **union**

```
val odd: RDD[Int] = sc.parallelize(List(1,3,5))
```

```
val even: RDD[Int] = sc.parallelize(List(2,4,6))
```

```
val all: RDD[Int] = odd.union(even)
```



RDD - Actions

- It launches the RDD evaluation,
 - returning some data to the driver
 - or persisting to some external storage system
- It's evaluated partially on each worker node, and results are merged in the Driver
- There are mechanisms to make cheaper computing several times the same RDD (like `persist()` or `cache()`)



RDD - Actions

- Examples:

- **count**

```
val myRdd: Rdd[Int] = sc.parallelize(List(1,2,3))
```

```
val size: Int = myRdd.count
```

Counting implies processing whole RDD

- **take**

```
val myRdd: Rdd[Int] = sc.parallelize(List(1,2,3))
```

```
val List(first,second) = myRdd.take(2).toList
```

- **collect**

```
val myRdd: Rdd[Int] = sc.parallelize(List(1,2,3))
```

```
val data: Array[Int] = myRdd.collect
```

Beware! Executing 'collect' on big collections might end into a memory leak



Demo1 : Most retweeted

- [Most Retweeted example](#)
- Bunch of tweets (1000 json records)
- Find out which is the most retweeted tweet.



Key-Value RDDs

- Particular case of $\text{RDD}[\text{T}]$ where $\text{T} = (\text{U}, \text{V})$
- It allows grouping, combining, aggregating values by some key.
- In Scala it's only needed to
`import org.apache.spark.SparkContext._`
- In Java, it's mandatory to use `PairRDD` class.



Key-Value RDDs

- keyBy: Generates a PairRDD

```
val myRDD: RDD[Int] =  
  sc.parallelize(List(1, 2, 3, 4, 5))
```

```
val kvRDD: RDD[(String, Int)] =  
  myRDD.keyBy(  
    n => if (n%2==0) "even" else "odd")
```

"odd" -> 1, "even" -> 2, "odd" -> 3, "even" -> 4, "odd" -> 5



Key-Value RDDs

- keys: Gets keys from PairRDD

```
val myRDD: RDD[(String, Int)] =  
    sc.parallelize("odd" -> 1, "even" -> 2, "odd" -> 3)  
val keysRDD: RDD[String] = myRDD.keys
```

"odd", "even", "odd"

- values: Gets values from PairRDD



Key-Value RDDs

- `mapValues`: map PairRDD values, omitting keys.

```
val myRDD: RDD[(String, Int)] =  
  sc.parallelize("odd" -> 1, "even" -> 2, "odd" -> 3)  
  
val mapRDD: RDD[(String, String)] =  
  myRDD.mapValues(_.toString)
```

"odd" -> "1", "even" -> "2", "odd" -> "3"

- `flatMapValues`: flatMap PairRDD values



Key-Value RDDs

- join: Return a new RDD with both RDD joined by key.

```
val a = sc.parallelize(List("dog", "salmon", "salmon", "rat", "elephant"), 3)
val b = a.keyBy(_.length)
val c = sc.parallelize(List("dog", "cat", "gnu", "salmon", "rabbit", "turkey", "wolf",
  "bear", "bee"), 3)
val d = c.keyBy(_.length)
b.join(d).collect

res0: Array[(Int, (String, String))] = Array((6,(salmon,salmon)), (6,(salmon,rabbit)),
(6,(salmon,turkey)), (6,(salmon,salmon)), (6,(salmon,rabbit)), (6,(salmon,turkey)), (3,
(dog,dog)), (3,(dog,cat)), (3,(dog,gnu)), (3,(dog,bee)), (3,(rat,dog)), (3,(rat,cat)),
(3,(rat,gnu)), (3,(rat,bee)))
```



Key-Value RDDs - combineByKey

- `combineByKey`:

```
def combineByKey [C] (  
    createCombiner: V => C,  
    mergeValue: (C, V) => C,  
    mergeCombiners: (C, C) => C) : RDD[(K, C)]
```

- Think of it as *something-like-but-not* a `foldLeft` over each partition λ



Key-Value RDDs - combineByKey

- It's composed by:
 - `createCombiner (V => C)` : Sets the way to mutate initial RDD[V] data into new data type used for aggregating values (C). This will be called Combinator.
 - `mergeValue ((C, V) => C)` : Defines how to aggregate initial V values to our Combiner type C, returning a new combiner type C.
 - `mergeCombiners ((C, C) => C)` : Defines how to merge two combiners into a new one.



Key-Value RDDs - combineByKey

- Example:

```
val a = sc.parallelize(List("dog", "cat", "gnu", "salmon", "rabbit", "turkey", "wolf", "bear", "bee"), 3)  
  
val b = sc.parallelize(List(1,1,2,2,2,1,2,2,2), 3)  
  
val c = b.zip(a)  
  
val d = c.combineByKey(List(_), (x:List[String], y:String) => y :: x, (x:List[String], y:List[String]) => x :::: y)  
  
d.collect
```

res16: Array[(Int, List[String])] = Array((1,List(cat, dog, turkey)), (2,List(gnu, rabbit, salmon, bee, bear, wolf)))



Key-Value RDDs - aggregateByKey

- **aggregateByKey** : Aggregate the values of each key, using given combine functions and a neutral “zero value”.
- *Beware!* Zero value is evaluated in each partition.

```
def aggregateByKey[U:ClassTag] (zeroValue: U) (seqOp: (U,V) => U, combOp: (U,U) => U): RDD[(K,U)] =  
  combineByKey(  
    (v: V) => seqOp(zeroValue,v),  
    seqOp,  
    combOp)
```



Key-Value RDDs - groupByKey

- **groupByKey**: Group all values with same key into a Iterable of values

```
def groupByKey(): RDD[(Key, Iterable[Value])] =  
  combineByKey[List[Value]](  
    (v: V) => List(v),  
    (list: List[V], v: V) => list += v,  
    (l1: List[V], l2: List[V]) => l1 ++ l2)
```

Note: GroupByKey actually uses CompactBuffer instead of list.



Modules

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Apache Spark





SparkSQL

- Shark successor (Berkeley)
 - Alternative to Hive + Hadoop
- Query language
 - HiveQL
 - Future: Support full SQL92
- SQLContext

```
val sqlContext = new SQLContext(sparkContext)
```

```
val hiveContext = new HiveContext(sparkContext)
```



SparkSQL

- Different datasources
 - JSON
 - Parquet
 - CSV
- New implementations
 - Cassandra
 - ElasticSearch
 - MongoDB
- Unified API in Spark 1.3.0

```
val students: DataFrame = sqlContext.load(  
  "students", "org.apache.sql.parquet", Map(...))
```



SparkSQL - DataFrame

- DataFrame = RDD[org.apache.spark.sql.Row] + Schema
- A Row holds both column values and their types.
- This allows unifying multiple datasources with the same API.
i.e, we could join two tables, one declared on Mongo and another on ElasticSearch.



Demo2 : Most retweeted(SparkSQL)

- [Most Retweeted Example \(SparkSQL\)](#)
- Same bunch of tweets (1000 json records)
- Find out which is the most retweeted tweet using SparkSQL



Who uses it?

amazon

databricks

YAHOO!



JPL

National Aeronautics and Space
Administration
Jet Propulsion Laboratory
California Institute of Technology

ebay™



Nokia
Networks



cloudera®



STRATIO



TREND
MICRO™



Let's talk about...

- Spark streaming
- Machine learning (MLlib)
- GraphX vs Neo4j
- Spark vs Hadoop
- ?







Introduction to



Javier Santos

May, 11 - 2015